

Machine Learning & Pattern Recognition

Xin Xin(辛鑫)

xinxin@sdu.edu.cn

<https://xinxin-me.github.io/>

Linear Regression

Linear Regression

age	23 years
annual salary	NTD 1,000,000
year in job	0.5 year
current debt	200,000

Training dataset: $\mathcal{D} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_m, y_m)\};$

Features of the i -th customer: $\mathbf{x}_i = (x_{i1} \ x_{i2} \ \dots \ x_{id})^T$; (Column vector)

The **ground truth** of the credit limit for the i -th customer: $y_i \in \mathbb{R}$.

Linear regression: $h(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i + b = \sum_{j=1}^d w_j x_{ij} + b$, where $\mathbf{w} = (w_1 \ w_2 \ \dots \ w_d)^T \in \mathbb{R}^d$

For simplicity, the bias b can be merged into the weight \mathbf{w} :

$$h(\mathbf{x}_i) = \hat{\mathbf{w}}^T \hat{\mathbf{x}}_i \quad \begin{aligned} \hat{\mathbf{w}} &= (b; \mathbf{w}) = (b \ w_1 \ w_2 \ \dots \ w_d) \in \mathbb{R}^{d+1} \\ \hat{\mathbf{x}}_i &= (1; x_{i1}; x_{i2}; \dots; x_{id}) \in \mathbb{R}^{d+1} \end{aligned}$$

Linear Regression



To-be-learned parameter

Linear regression hypothesis: $h(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i = \sum_{j=0}^d w_j x_{ij}$, $x_{i0} = 1$

Linear regression: find lines/hyperplanes with small **residuals**.

Popular/historical squared error measure:

$$L(h(\mathbf{x}), y) = (\hat{y} - y)^2$$



Empirical Error

We prefer to minimize the objective function where the expectation is taken across the **data generating distribution** p_{data} rather than just over the finite training set:

$$J^*(\boldsymbol{\theta}) = \mathbb{E}_{(\mathbf{x}, y) \sim p_{data}} L(h(\mathbf{x}, \boldsymbol{\theta}), y)$$

However, in most cases, we do not know p_{data} but only have a training set of samples. One simplest way to convert the machine learning problem back into an optimization problem is to minimize the expected loss on the training set.

$$J(\boldsymbol{\theta}) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{P}_{data}} L(h(\mathbf{x}, \boldsymbol{\theta}), y)$$

Replacing the **true** distribution $p_{data}(\mathbf{x}, y)$ with the **empirical** distribution $\hat{P}_{data}(\mathbf{x}, y)$ defined by the training set.

Linear Regression

Popular/historical error measure:

$$\text{squared error } L(h(\mathbf{x}), y) = (\hat{y} - y)^2$$

$$E(\mathbf{w}) = \sum_{i=1}^m \frac{h(\mathbf{x}_i) - y_i}{\mathbf{w}^T \mathbf{x}_i}^2$$

Next: How to minimize $E(\mathbf{w})$?

Matrix Form of $E(\mathbf{w})$

$$loss = \sum_{i=1}^m (\mathbf{w}^T \mathbf{x}_i - y_i)^2 + \lambda \|\mathbf{w}\|^2, \quad \mathbf{w} = (w_0, w_1, \dots, w_d)^T$$

$$E(\mathbf{w}) = \sum_{i=1}^m (h(\mathbf{x}_i) - y_i)^2 = \sum_{i=1}^m (\mathbf{w}^T \mathbf{x}_i - y_i)^2 = \sum_{i=1}^m (\mathbf{x}_i^T \mathbf{w} - y_i)^2$$

$$= \left\| \begin{bmatrix} \mathbf{x}_1^T \mathbf{w} - y_1 \\ \mathbf{x}_2^T \mathbf{w} - y_2 \\ \vdots \\ \mathbf{x}_m^T \mathbf{w} - y_m \end{bmatrix} \right\|^2 = \left\| \begin{bmatrix} - & - & \mathbf{x}_1^T & - & - \\ - & - & \mathbf{x}_2^T & - & - \\ & & \vdots & & \\ - & - & \mathbf{x}_m^T & - & - \end{bmatrix} \mathbf{w} - \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix} \right\|^2$$

$$= \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \quad l_2\text{-norm } \|\mathbf{x}\|_2 = \sqrt{x_1^2 + x_2^2 + \dots + x_d^2}$$


The subscript '2' is usually omitted.

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1d} \\ 1 & x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{m1} & x_{m2} & \cdots & x_{md} \end{pmatrix} \in \mathbb{R}^{m \times (d+1)}, \quad \mathbf{w} = \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{pmatrix} \in \mathbb{R}^{d+1}, \quad \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \in \mathbb{R}^m$$

Matrix Form of $E(\mathbf{w})$

A continuous, twice differentiable function of several variables is convex on a convex set if and only if its Hessian matrix is positive semidefinite on the interior of the convex set.

$$\min E(\mathbf{w}) = \min \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2$$



- $E(\mathbf{w})$: continuous, differentiable, convex
- Necessary condition of 'best' \mathbf{w} .

$$\nabla E(\mathbf{w}) = \begin{bmatrix} \frac{\partial E}{\partial w_0}(\mathbf{w}) \\ \frac{\partial E}{\partial w_1}(\mathbf{w}) \\ \vdots \\ \frac{\partial E}{\partial w_d}(\mathbf{w}) \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \quad \checkmark \text{ Not possible to 'roll down'}$$

Task: find the \mathbf{w}^* such that $\nabla E(\mathbf{w}^*) = 0$

The Gradient $\nabla E(\mathbf{w})$

$$\min_{\mathbf{w}} E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = \underbrace{\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}}_A - 2 \underbrace{\mathbf{w}^T \mathbf{X}^T \mathbf{y}}_b + \underbrace{\mathbf{y}^T \mathbf{y}}_c$$

One w only

$$E(w) = (aw^2 - 2bw + c)$$

$$\nabla E(w) = 2aw - 2b$$

Vector w

$$E(\mathbf{w}) = (\mathbf{w}^T A \mathbf{w} - 2\mathbf{w}^T \mathbf{b} + c)$$

$$\nabla E(\mathbf{w}) = ?$$

Derivatives

■ **scalar – scalar:** e.g., $\frac{d}{dx} x^2 = 2x$

■ **scalar-vector:** e.g., $f(\mathbf{x})$ is a scalar function of vector \mathbf{x}

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_d \end{bmatrix} \quad \frac{df}{d\mathbf{x}} = \begin{bmatrix} \frac{\sigma f}{\sigma x_1} \\ \vdots \\ \frac{\sigma f}{\sigma x_d} \end{bmatrix}$$

■ **scalar-matrix:** $f(\mathbf{A})$ is a scalar function and $m \times n$ matrix \mathbf{A}

$$\frac{df}{d\mathbf{A}} = \begin{bmatrix} \frac{\sigma f}{\sigma a_{11}} & \dots & \frac{\sigma f}{\sigma a_{1d}} \\ \vdots & \ddots & \vdots \\ \frac{\sigma f}{\sigma a_{m1}} & \dots & \frac{\sigma f}{\sigma a_{mn}} \end{bmatrix}$$

w.r.t

	Differentiate		
	scalar	vector	matrix
scalar	scalar	vector	matrix
vector	vector	matrix	
matrix	matrix		

Matrix Calculus

- Numerator layout: lay out according to \mathbf{y} and \mathbf{x}^T . (Jacobian formulation)
- Denominator layout: lay out according to \mathbf{y}^T and \mathbf{x} . (Hessian formulation)

Numerator layout:

分子布局

$$\frac{\partial y}{\partial \mathbf{x}} = \left[\frac{\partial y}{\partial x_1} \frac{\partial y}{\partial x_2} \cdots \frac{\partial y}{\partial x_n} \right]$$

$$\frac{\partial \mathbf{y}}{\partial x} = \begin{bmatrix} \frac{\partial y_1}{\partial x} \\ \frac{\partial y_2}{\partial x} \\ \vdots \\ \frac{\partial y_n}{\partial x} \end{bmatrix}$$

Denominator layout:

分母布局

$$\frac{\partial y}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial y}{\partial x_1} \\ \frac{\partial y}{\partial x_2} \\ \vdots \\ \frac{\partial y}{\partial x_n} \end{bmatrix}$$

$$\frac{\partial \mathbf{y}}{\partial x} = \left[\frac{\partial y_1}{\partial x} \frac{\partial y_2}{\partial x} \cdots \frac{\partial y_n}{\partial x} \right]$$

Commonly Used Derivatives

- $\frac{d}{dx}(Ax) = A^T$

- $\frac{dx}{dx} = I$

- $\frac{dy^T x}{dx} = \frac{dx^T y}{dx} = y$

- $\frac{d}{dx}(x^T Ax) = \begin{cases} (A + A^T)x & \text{If } \mathbf{A} \text{ square} \\ 2Ax & \text{If } \mathbf{A} \text{ symmetric} \end{cases}$

The Gradient $\nabla E(\mathbf{w})$

$$\min_{\mathbf{w}} E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y}) = \underbrace{\mathbf{w}^T \mathbf{X}^T \mathbf{X} \mathbf{w}}_A - 2 \underbrace{\mathbf{w}^T \mathbf{X}^T \mathbf{y}}_b + \underbrace{\mathbf{y}^T \mathbf{y}}_c$$

One w only

$$E(w) = (aw^2 - 2bw + c)$$

$$\nabla E(w) = 2aw - 2b$$

Vector w

$$E(\mathbf{w}) = (\mathbf{w}^T \mathbf{A} \mathbf{w} - 2\mathbf{w}^T \mathbf{b} + c)$$

$$\nabla E(\mathbf{w}) = 2\mathbf{A}\mathbf{w} - 2\mathbf{b}$$

$$\nabla E(\mathbf{w}) = 2(\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y})$$

Optimal Linear Regression Weights

Task: find \mathbf{w}^* such that $\nabla E(\mathbf{w}^*) = 2(\mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y}) = \mathbf{0}$

Invertible/positive definite $\mathbf{X}^T \mathbf{X}$

- Unique solution

$$\mathbf{w}^* = (\underbrace{\mathbf{X}^T \mathbf{X}})^{-1} \mathbf{X}^T \mathbf{y}$$

pseudo-inverse \mathbf{X}^\dagger

Note: $\mathbf{X}^\dagger \mathbf{X} = \mathbf{I}$, but $\mathbf{X} \mathbf{X}^\dagger \neq \mathbf{I}$

If \mathbf{X} is square and invertible, $\mathbf{X}^\dagger = \mathbf{X}^{-1}$.

Singular $\mathbf{X}^T \mathbf{X}$

- Define \mathbf{X}^\dagger in other ways (e.g., SVD).
- Add regularization
 - E.g., l_2 norm $\lambda > 0$

$$\min E(\mathbf{w}) = \min \|\mathbf{X} \mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2$$

$$\nabla E(\mathbf{w}^*) = 2(\mathbf{X}^T \mathbf{X} \mathbf{w} + \lambda \mathbf{w} - \mathbf{X}^T \mathbf{y}) = \mathbf{0}$$

$$\underline{(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I}) \mathbf{w} = \mathbf{X}^T \mathbf{y}}$$

Invertible?

Linear Regression Algorithm

1. From \mathcal{D} , construct input matrix \mathbf{X} and output vector \mathbf{y} by

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1d} \\ 1 & x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{m1} & x_{m2} & \cdots & x_{md} \end{pmatrix} \in \mathbb{R}^{m \times (d+1)}, \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \in \mathbb{R}^m$$

2. Calculate pseudo-inverse

$$\mathbf{X}^\dagger \in \mathbb{R}^{(d+1) \times m}$$

3. Return $\mathbf{w}^* = \mathbf{X}^\dagger \mathbf{y} \in \mathbb{R}^{(d+1)}$

Simple and efficient (?) with **good \mathbf{X}^\dagger**

Logistic Regression

Heart Attack Prediction Problem

age	40 years
gender	male
blood pressure	130/85
cholesterol level	240
weight	70

heart disease? **yes**

Binary classification:

Ideal $f(\mathbf{x}) = \text{sign}(p(+1|\mathbf{x}) - 0.5) \in \{-1, +1\}$

Heart Attack Prediction Problem

age	40 years
gender	male
blood pressure	130/85
cholesterol level	240
weight	70

heart attack? **80% risk**

'Soft' Binary classification:

$$f(x) = p(+1|x) \in [0,1]$$

Soft Binary classification:

Target function $f(x) = p(+1|x) \in [0,1]$

Ideal data

$$\begin{pmatrix} \mathbf{x}_1, y'_1 = 0.9 = P(+1|\mathbf{x}_1) \\ \mathbf{x}_2, y'_2 = 0.2 = P(+1|\mathbf{x}_2) \\ \vdots \\ \mathbf{x}_N, y'_N = 0.6 = P(+1|\mathbf{x}_N) \end{pmatrix}$$

Actual data

$$\begin{pmatrix} \mathbf{x}_1, y_1 = \circ \sim P(y|\mathbf{x}_1) \\ \mathbf{x}_2, y_2 = \times \sim P(y|\mathbf{x}_2) \\ \vdots \\ \mathbf{x}_N, y_N = \times \sim P(y|\mathbf{x}_N) \end{pmatrix}$$

Same data as hard binary classification, different **target function**

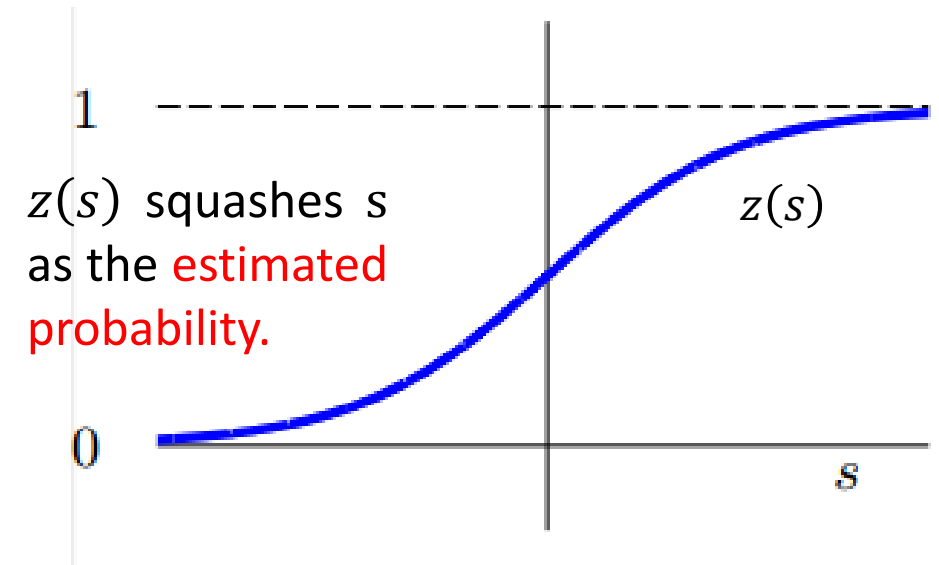
Logistic Hypothesis

age	40 years
gender	male
blood pressure	130/85
cholesterol level	240

Let $\mathbf{x}_i = (x_{i0}, x_{i1}, x_{i2}, \dots, x_{id})$ be the features of the patient, calculate a weighted 'risk score':

$$s = \sum_{j=0}^d w_j x_{ij} = \mathbf{w}^T \mathbf{x}_i,$$


Convert the score to estimated probability by logistic function $z(s)$.



$$\text{Logistic hypothesis: } h(\mathbf{x}_i) = z(\mathbf{w}^T \mathbf{x}_i)$$

Logistic Function

$$z(s) = \frac{e^s}{1 + e^s} = \frac{1}{1 + e^{-s}}$$



smooth, monotonic,
sigmoid function of s

Bound	$z(s) \in [0,1]$	$z(-\infty) = 0$	$z(0) = 0.5$	$z(\infty) = 1$
Symmetric	$1 - z(s) = z(-s)$			
Gradient	$z'(s) = z(s)(1 - z(s))$			

Logistic regression use $h(\mathbf{x}) = z(\mathbf{w}^T \mathbf{x})$ to approximate the target $f(\mathbf{x}) = p(+1|\mathbf{x})$

Exercise

Logistic Regression and Binary Classification

Consider any logistic hypothesis $h(\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$ that approximates $P(y|\mathbf{x})$. 'Convert' $h(\mathbf{x})$ to a binary classification prediction by taking $\text{sign}(h(\mathbf{x}) - \frac{1}{2})$. What is the equivalent formula for the binary classification prediction?

- 1 $\text{sign}(\mathbf{w}^T \mathbf{x} - \frac{1}{2})$
- 2 $\text{sign}(\mathbf{w}^T \mathbf{x})$
- 3 $\text{sign}(\mathbf{w}^T \mathbf{x} + \frac{1}{2})$
- 4 none of the above

Exercise

Logistic Regression and Binary Classification

Consider any logistic hypothesis $h(\mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$ that approximates $P(y|\mathbf{x})$. 'Convert' $h(\mathbf{x})$ to a binary classification prediction by taking $\text{sign}(h(\mathbf{x}) - \frac{1}{2})$. What is the equivalent formula for the binary classification prediction?

- 1 $\text{sign}(\mathbf{w}^T \mathbf{x} - \frac{1}{2})$
- 2 $\text{sign}(\mathbf{w}^T \mathbf{x})$
- 3 $\text{sign}(\mathbf{w}^T \mathbf{x} + \frac{1}{2})$
- 4 none of the above


Reference Answer: 2

When $\mathbf{w}^T \mathbf{x} = 0$, $h(\mathbf{x})$ is exactly $\frac{1}{2}$. So thresholding $h(\mathbf{x})$ at $\frac{1}{2}$ is the same as thresholding $(\mathbf{w}^T \mathbf{x})$ at 0.

Linear Models

linear regression


$$h(\mathbf{x}) = s$$



friendly err = squared
(easy to minimize)

logistic regression

$$h(\mathbf{x}) = z(s)$$



err = ?

How to define the cost (error) function for logistic regression?

Maximum-Likelihood Estimation


Given a dataset $\mathcal{D} = \{x_1, x_2, \dots, x_n\}$, where the n samples are drawn **independently** from **identical** distribution $p(x|\theta)$, estimate parameters θ .

ML estimate parameters θ maximizes $p(\mathcal{D}|\theta)$

\mathcal{D} is an i.i.d set

$$\hat{\theta} = \arg \max_{\theta} p(\mathcal{D}|\theta)$$

$$p(\mathcal{D}|\theta) = \prod_{k=1}^n p(x_k|\theta)$$



Logistic Regression-- $y \in \{0,1\}$

Consider $\mathcal{D} = \{(\mathbf{x}_1, +), (\mathbf{x}_2, -), \dots, (\mathbf{x}_m, -)\}$

Likelihood that h generates \mathcal{D}

$$\begin{aligned} & p(\mathbf{x}_1)h(\mathbf{x}_1) \\ & p(\mathbf{x}_2)(1 - h(\mathbf{x}_2)) \\ & \vdots \\ & p(\mathbf{x}_m)(1 - h(\mathbf{x}_m)) \end{aligned}$$

- Target function:
 $f(x) = p(+1|x)$
- If $h \approx f$, then likelihood
 $(h) \approx$ that using (f)

Likelihood of Logistic Regression

Goal: $\arg \max_h \text{likelihood}(h)$ $\text{likelihood}(h) = \prod_{i=1}^m p(\mathbf{x}_i) p(y|\mathbf{x}_i)$

Consider $\mathcal{D} = \{(\mathbf{x}_1, +), (\mathbf{x}_2, -), \dots, (\mathbf{x}_m, -)\}$

$$\begin{aligned} \text{likelihood}(h) &= \prod_{i=1}^m p(\mathbf{x}_i) p(y_i|\mathbf{x}_i) \\ &= p(\mathbf{x}_1) h(\mathbf{x}_1) p(\mathbf{x}_2) (1 - h(\mathbf{x}_2)) \cdots p(\mathbf{x}_m) (1 - h(\mathbf{x}_m)) \end{aligned}$$

Likelihood of Logistic Regression

Goal: $\arg \max_h \text{likelihood}(h)$ $\text{likelihood}(h) = \prod_{i=1}^m p(\mathbf{x}_i) p(y|\mathbf{x}_i)$

Consider $\mathcal{D} = \{(\mathbf{x}_1, +), (\mathbf{x}_2, -), \dots, (\mathbf{x}_m, -)\}$

$$\begin{aligned} \text{likelihood}(h) &= \prod_{i=1}^m p(\mathbf{x}_i) p(y_i|\mathbf{x}_i) \\ &= p(\mathbf{x}_1) h(\mathbf{x}_1) p(\mathbf{x}_2) (1 - h(\mathbf{x}_2)) \cdots p(\mathbf{x}_m) (1 - h(\mathbf{x}_m)) \end{aligned}$$

We remove all the $p(\mathbf{x}_i)$ which remains the same for all the hypothesis h .

Likelihood of Logistic Regression

$$\text{likelihood}(\mathbf{h}) = \prod_{i=1}^m p(\mathbf{x}_i) p(y_i | \mathbf{x}_i) \propto \prod_{i=1}^m p(y_i | \mathbf{x}_i)$$

$$p(y_i | \mathbf{x}_i) = \begin{cases} h(\mathbf{x}_i) & \text{for } y_i = 1 \\ 1 - h(\mathbf{x}_i) & \text{for } y_i = 0 \end{cases} \iff p(y_i | \mathbf{x}_i) = h(\mathbf{x}_i)^{y_i} (1 - h(\mathbf{x}_i))^{(1-y_i)}$$

Bernoulli distribution

$$\text{likelihood}(h) \propto \prod_{i=1}^m p(y_i | \mathbf{x}_i) = \prod_{i=1}^m h(\mathbf{x}_i)^{y_i} (1 - h(\mathbf{x}_i))^{(1-y_i)}$$

Log-Likelihood of Logistic Regression

Negative Log-likelihood

$$\min_{\mathbf{h}} E(\mathbf{h}) = \sum_{i=1}^m -(y_i \ln \mathbf{h}(\mathbf{x}_i) + (1 - y_i) \ln(1 - \mathbf{h}(\mathbf{x}_i)))$$

Cross-entropy loss

Cross-entropy

$$H(\mathbf{p}, \mathbf{q}) = - \sum_x \mathbf{p}(\mathbf{x}) \log(\mathbf{q}(\mathbf{x}))$$

$\mathbf{p} \in \{y, 1 - y\}$
 $\mathbf{q} \in \{\mathbf{h}(\mathbf{x}), 1 - \mathbf{h}(\mathbf{x})\}$

Negative Log-likelihood


$$\min_{\mathbf{w}} \sum_{i=1}^m \left[-y_i \ln \left(\frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}_i}} \right) - (1 - y_i) \ln \left(\frac{1}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} \right) \right]$$

$$\min_{\mathbf{w}} \sum_{i=1}^m \left[-y_i \mathbf{w}^T \mathbf{x}_i + \ln(1 + e^{\mathbf{w}^T \mathbf{x}_i}) \right]$$

Minimize $E(w)$

$$\min_w E(w) = \sum_{i=1}^m \left[-y_i w^T x_i + \ln(1 + e^{w^T x_i}) \right]$$

Cross-entropy loss



$E(w)$: continuous, differentiable,
twice-differentiable, **convex**
We want to find the valley

$$\nabla E(w) = 0$$

Matrix Calculus

$$\min_{\mathbf{w}} E(\mathbf{w}) = \sum_{i=1}^m \left[-y_i \mathbf{w}^T \mathbf{x}_i + \ln(1 + e^{\mathbf{w}^T \mathbf{x}_i}) \right]$$

Identities: scalar-by-vector $\frac{\partial y}{\partial \mathbf{x}} = \nabla_{\mathbf{x}} y$

Condition	Expression	Numerator layout, i.e. by \mathbf{x}^T ; result is row vector	Denominator layout, i.e. by \mathbf{x} ; result is column vector
a is not a function of \mathbf{x}	$\frac{\partial a}{\partial \mathbf{x}} =$	$\mathbf{0}^T$ [4]	$\mathbf{0}$ [4]
a is not a function of \mathbf{x} , $u = u(\mathbf{x})$	$\frac{\partial au}{\partial \mathbf{x}} =$		$a \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial(u+v)}{\partial \mathbf{x}} =$		$\frac{\partial u}{\partial \mathbf{x}} + \frac{\partial v}{\partial \mathbf{x}}$
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial uv}{\partial \mathbf{x}} =$		$u \frac{\partial v}{\partial \mathbf{x}} + v \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x})$	$\frac{\partial g(u)}{\partial \mathbf{x}} =$		$\frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x})$	$\frac{\partial f(g(u))}{\partial \mathbf{x}} =$		$\frac{\partial f(g)}{\partial g} \frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$
$\mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial(\mathbf{u} \cdot \mathbf{v})}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}^T \mathbf{v}}{\partial \mathbf{x}} =$	$\mathbf{u}^T \frac{\partial \mathbf{v}}{\partial \mathbf{x}} + \mathbf{v}^T \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$ • assumes numerator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{v} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u}$ • assumes denominator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$

Gradient $\nabla E(\mathbf{w})$

$$\nabla E(\mathbf{w}) = \sum_{i=1}^m \left[-y_i \mathbf{x}_i + \frac{e^{\mathbf{w}^T \mathbf{x}_i}}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} \mathbf{x}_i \right] = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i = 0$$

- $\nabla E(\mathbf{w})$ is a non-linear equation of \mathbf{w}
 - It is hard to derive the **closed form** solution. :-)

Gradient $\nabla E(\mathbf{w})$


$$\nabla E(\mathbf{w}) = \sum_{i=1}^m \left[-y_i \mathbf{x}_i + \frac{e^{\mathbf{w}^T \mathbf{x}_i}}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} \mathbf{x}_i \right] = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i = 0$$

- Apply the iterative optimization to the logistic regression.

Iterative Optimization

Optimization Methods

- Optimization: either minimize or maximize some function $f(x)$ by altering x .
- In most cases, optimization refers to the minimization of $f(x)$.

Maximization $f(x)$  **Minimization** $-f(x)$

- $f(x)$: objective function, cost function, loss function, error function.
- The value that minimize $f(x)$: $x^* = \arg \min f(x)$.

Optimization Methods

- **Deterministic Optimization**
 - The data for the given problem are known accurately.
- **Stochastic Optimization**
 - Refers to a collection of methods for minimizing or maximizing an objective function when randomness is present.

Deterministic Optimization

- First-order methods: methods that use only the **gradient**.
- Second-order methods: methods that also use the **Hessian** matrix.

$$\mathbf{H}(f)_{i,j} = \frac{\partial^2}{\partial x_i \partial x_j} f(\mathbf{x})$$

\mathbf{x} : multiple input dimensions.

Taylor Approximation

Expansion at x_0

$$f(x) = \frac{f(x_0)}{0!} + \frac{f'(x_0)}{1!}(x - x_0) + \frac{f''(x_0)}{2!}(x - x_0)^2 + \dots + \frac{f^{(n)}(x_0)}{n!}(x - x_0)^n + R_n(x)$$

Examples

$$e^x = 1 + \frac{1}{1!}x + \frac{1}{2!}x^2 + \frac{1}{3!}x^3 + o(x^3)$$

$$\ln(1+x) = x - \frac{1}{2}x^2 + \frac{1}{3}x^3 + o(x^3)$$

$$\sin x = x - \frac{1}{3!}x^3 + \frac{1}{5!}x^5 + o(x^5)$$

Gradient Descent [Cauchy 1847]

- Motivation: to **minimize** the local **first-order Taylor approximation** of f

$$\min_x f(x) \approx \min_x f(x_t) + \nabla f(x_t)^T (x - x_t)$$


- Update rule:

$$x_{t+1} = x_t - \eta_t \nabla f(x_t)$$

Where $\eta_t > 0$ is the step-size (learning rate).

Interpretation

- Reduce $f(x)$ by moving x in small steps with opposite sign of the derivative.
- Update rule:
$$x_{t+1} = x_t - \eta_t \nabla f(x_t)$$
- **Critical/stationary** points: Points where $f'(x) = 0$ 驻点



An illustration of gradient descent.

Interpretation

- At each iteration, consider the expansion


$$f(x) \approx \underbrace{f(x_t) + \nabla f(x_t)^T (x - x_t)}_{\text{Linear approximation of } f} + \underbrace{\frac{1}{2\eta_t} \|x - x_t\|^2}_{\text{Proximity term with weight } \frac{1}{2\eta_t}}$$

- Quadratic approximation, replacing usual $\nabla^2 f(x)$ by $\frac{1}{\eta_t} I$:

$$x_{t+1} = x_t - \eta_t \nabla f(x_t)$$

Interpretation

$$f(x) \approx f(x_t) + \nabla f(x_t)^T (x - x_t) + \frac{1}{2\eta_t} \|x - x_t\|^2$$




Blue point is x_t , red point is x_{t+1} .

Global VS Local Minimum


- Global minimum: a point that obtains the absolute lowest value of $f(x)$.
- Local minimum: a point where $f(x)$ is lower than at all neighboring points.
- Saddle points: some critical points are neither maxima or minima. 鞍点

Types of critical points


Minimum



Maximum




Saddle points




Global VS Local Minimum

- Global minimum: a point that obtains the absolute lowest value of $f(x)$.
- Local minimum: a point where $f(x)$ is higher than at all neighboring points.
- Saddle points: some critical points are neither maxima or minima. 鞍点




Different Starting Points

- Gradient Descent with different starting points are illustrated in different colors.



(a) Convex function



(b) Non-convex function

- (a): Strictly convex function: Converge to the global optimum.
- (b): Non-convex function: Different paths may end up at different local optima.

Gradient Descent [Cauchy 1847]


$$x_{t+1} = x_t - \eta_t \nabla f(x_t)$$

- Gradient Descent requires a step size η controlling the amount of gradient updated to the current point at each iteration.
- It is naïve to set $\eta_t = \eta$ for all iterations.


How to choose step sizes?

Fixed Step Sizes

Considering $f(x) = (10x_1^2 + x_2^2)/2$



3D Plot




Contour Plot

Fixed Step Sizes

Considering $f(x) = (10x_1^2 + x_2^2)/2$

If η is too big, can lead to divergence.

- The learning function oscillates away from the optimal point.
- As shown, it oscillates after 8 steps.




Fixed Step Sizes

Considering $f(x) = (10x_1^2 + x_2^2)/2$

If η is too small, takes longer time for the function to converge.


- As shown, GD after 100 steps.



Fixed Step Sizes

Considering $f(x) = (10x_1^2 + x_2^2)/2$


Same example, gradient descent after 40 appropriately sized steps.




Fixed Step Sizes

Considering $f(x) = x^2/2$

Big learning rate



Small learning rate



Deterministic Optimization

- First-order methods: methods that use only the **gradient**.
- **Second-order methods: methods that also use the Hessian matrix.**

Suppose $f: \mathbb{R}^n \rightarrow \mathbb{R}$ is a function taking as input a vector $x \in \mathbb{R}^n$ and outputting a scalar $f(x) \in \mathbb{R}$; if all second partial derivatives of f exist and are continuous over the domain of the function, then the Hessian matrix \mathbf{H} of f is a square $n \times n$ matrix, usually defined as follows.

$$H = \nabla^2 f = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 x_n} \\ \frac{\partial^2 f}{\partial x_2 x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n x_1} & \frac{\partial^2 f}{\partial x_n x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}, \text{ or } H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j}$$

Newton's Methods

- Motivation: to minimize the local **second-order Taylor** approximation of f .

$$\min_{\mathbf{x}} f(\mathbf{x}) \approx \min_{\mathbf{x}} f(\mathbf{x}_t) + \nabla f(\mathbf{x}_t)^T (\mathbf{x} - \mathbf{x}_t) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_t)^T \nabla^2 f(\mathbf{x}_t) (\mathbf{x} - \mathbf{x}_t)$$

- Take the derivative of \mathbf{x} on both side, we have,

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = \nabla f(\mathbf{x}_t) + \nabla^2 f(\mathbf{x}_t) (\mathbf{x} - \mathbf{x}_t) = \mathbf{0}$$

- Update rule: suppose $\nabla^2 f(\mathbf{x}_t)$ is positive definite,

$$\mathbf{x} = \mathbf{x}_t - [\nabla^2 f(\mathbf{x}_t)]^{-1} \nabla f(\mathbf{x}_t)$$

Newton's Methods

- Motivation: to minimize the local **second-order Taylor** approximation of f .

$$\min_x f(x) \approx \min_x f(x_t) + f'(x_t)(x - x_t) + \frac{1}{2}f''(x_t)(x - x_t)^2$$

- Take the derivative of x on both side, we have,

$$f'(x) = f'(x_t) + f''(x_t)(x - x_t) = 0$$


- Update rule: suppose $f''(x_t) \neq 0$,

$$x = x_t - \frac{f'(x_t)}{f''(x_t)}$$

Newton's Methods

- In numerical analysis, Newton's Methods is to find successively better approximations to the roots of a real-valued function, (i.e., $f(z) = 0$).

$$z = z_t - \frac{f(z_t)}{f'(z_t)}$$



- In optimization, we want to find the stationary point $f'(x_t) = 0$, i.e.,

$$x = x_t - \frac{f'(x_t)}{f''(x_t)}$$

Newton's Methods

- **Advantage:**

- More **accurate** local approximation of the objective,
- The convergence is much **faster**.

- **Disadvantage:**

- Need to compute the **second derivatives**
- Need to compute the **inverse** of Hessian (time/storage consuming)

Go back to logistic regression

Gradient $\nabla E(\mathbf{w})$

$$\nabla E(\mathbf{w}) = \sum_{i=1}^m \left[-y_i \mathbf{x}_i + \frac{e^{\mathbf{w}^T \mathbf{x}_i}}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} \mathbf{x}_i \right] = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i = \mathbf{X}^T (\hat{\mathbf{y}} - \mathbf{y})$$

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1d} \\ 1 & x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m1} & x_{m2} & \cdots & x_{md} \end{pmatrix} = \begin{pmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_m^T \end{pmatrix} \in \mathbb{R}^{m \times (d+1)}, \hat{\mathbf{y}} = \begin{pmatrix} z(\mathbf{w}^T \mathbf{x}_1) \\ z(\mathbf{w}^T \mathbf{x}_2) \\ \vdots \\ z(\mathbf{w}^T \mathbf{x}_m) \end{pmatrix} \in \mathbb{R}^m, \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \in \mathbb{R}^m$$

- Apply the Newton's method to the logistic regression,

$$\mathbf{x} = \mathbf{x}_t - [\nabla^2 f(\mathbf{x}_t)]^{-1} \nabla f(\mathbf{x}_t) \quad \longrightarrow \quad \mathbf{w} = \mathbf{w}_t - \mathbf{H}(\mathbf{w}_t)^{-1} \nabla E(\mathbf{w}_t)$$

- Need to solve,

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}} = ?$$

Gradient $\nabla E(\mathbf{w})$

$$\nabla E(\mathbf{w}) = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i$$

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}}$$

$$\mathbf{H} = \sum_{i=1}^m \frac{\nabla \{z(\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i\}}{\nabla \mathbf{w}}$$

Identities: vector-by-vector $\frac{\partial y}{\partial \mathbf{x}}$

Condition	Expression	Numerator layout, i.e. by y and \mathbf{x}^T	Denominator layout, i.e. by y^T and \mathbf{x}
\mathbf{a} is not a function of \mathbf{x}	$\frac{\partial \mathbf{a}}{\partial \mathbf{x}} =$	$\mathbf{0}$	
	$\frac{\partial \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{I}	
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{A} \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{A}	\mathbf{A}^T
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{x}^T \mathbf{A}}{\partial \mathbf{x}} =$	\mathbf{A}^T	\mathbf{A}
a is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	
$a = a(\mathbf{x})$, $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \mathbf{u} \frac{\partial a}{\partial \mathbf{x}}$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial a}{\partial \mathbf{x}} \mathbf{u}^T$
\mathbf{A} is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{A} \mathbf{u}}{\partial \mathbf{x}} =$	$\mathbf{A} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{A}^T$
$\mathbf{u} = \mathbf{u}(\mathbf{x})$, $\mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial (\mathbf{u} + \mathbf{v})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$	
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}}$
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{f}(\mathbf{g}(\mathbf{u}))}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}}$

Gradient $\nabla E(\mathbf{w})$

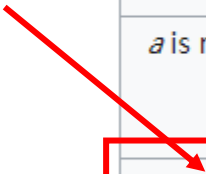
$$\nabla E(\mathbf{w}) = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i$$

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}}$$

$$\mathbf{H} = \sum_{i=1}^m \frac{\nabla \{z(\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i\}}{\nabla \mathbf{w}}$$

$\frac{\nabla z(\mathbf{w}^T \mathbf{x}_i)}{\nabla \mathbf{w}}$ is a scalar –by–vector problem.

$a: z(\mathbf{w}^T \mathbf{x}_i)$
 $\mathbf{u}(\mathbf{w}): \mathbf{x}_i$



Identities: vector-by-vector $\frac{\partial y}{\partial \mathbf{x}}$

Condition	Expression	Numerator layout, i.e. by y and \mathbf{x}^T	Denominator layout, i.e. by y^T and \mathbf{x}
a is not a function of \mathbf{x}	$\frac{\partial a}{\partial \mathbf{x}} =$	$\mathbf{0}$	
	$\frac{\partial \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{I}	
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{A} \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{A}	\mathbf{A}^T
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{x}^T \mathbf{A}}{\partial \mathbf{x}} =$	\mathbf{A}^T	\mathbf{A}
a is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	
$a = a(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \mathbf{u} \frac{\partial a}{\partial \mathbf{x}}$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial a}{\partial \mathbf{x}} \mathbf{u}^T$
\mathbf{A} is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{A} \mathbf{u}}{\partial \mathbf{x}} =$	$\mathbf{A} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{A}^T$
$\mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial (\mathbf{u} + \mathbf{v})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$	
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}}$
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{f}(\mathbf{g}(\mathbf{u}))}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}}$

Gradient $\nabla E(w)$

$\frac{\nabla_Z(w^T x_i)}{\nabla w}$ is a scalar –by-vector problem

Identities: scalar-by-vector $\frac{\partial y}{\partial \mathbf{x}} = \nabla_{\mathbf{x}} y$

Condition	Expression	Numerator layout, i.e. by \mathbf{x}^T ; result is row vector	Denominator layout, i.e. by \mathbf{x} ; result is column vector
a is not a function of \mathbf{x}	$\frac{\partial a}{\partial \mathbf{x}} =$	$\mathbf{0}^T$ [4]	$\mathbf{0}$ [4]
a is not a function of \mathbf{x} , $u = u(\mathbf{x})$	$\frac{\partial au}{\partial \mathbf{x}} =$	$a \frac{\partial u}{\partial \mathbf{x}}$	
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial(u+v)}{\partial \mathbf{x}} =$	$\frac{\partial u}{\partial \mathbf{x}} + \frac{\partial v}{\partial \mathbf{x}}$	
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial uv}{\partial \mathbf{x}} =$	$u \frac{\partial v}{\partial \mathbf{x}} + v \frac{\partial u}{\partial \mathbf{x}}$	
$u = u(\mathbf{x})$	$\frac{\partial g(u)}{\partial \mathbf{x}} =$	$\frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$	
$u = u(\mathbf{x})$	$\frac{\partial f(g(u))}{\partial \mathbf{x}} =$	$\frac{\partial f(g)}{\partial g} \frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$	
$\mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial(\mathbf{u} \cdot \mathbf{v})}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}^T \mathbf{v}}{\partial \mathbf{x}} =$	$\mathbf{u}^T \frac{\partial \mathbf{v}}{\partial \mathbf{x}} + \mathbf{v}^T \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$ <ul style="list-style-type: none"> assumes numerator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$ 	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{v} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u}$ <ul style="list-style-type: none"> assumes denominator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$

Gradient $\nabla E(w)$

$\frac{\nabla z(w^T x_i)}{\nabla w}$ is a scalar –by–vector problem $u: w^T x_i$ $z: g$ $\frac{\nabla z(w^T x_i)}{\nabla w} = z(w^T x_i) z(-w^T x_i) x_i$

Identities: scalar-by-vector $\frac{\partial y}{\partial \mathbf{x}} = \nabla_{\mathbf{x}} y$

Condition	Expression	Numerator layout, i.e. by \mathbf{x}^T ; result is row vector	Denominator layout, i.e. by \mathbf{x} ; result is column vector
a is not a function of \mathbf{x}	$\frac{\partial a}{\partial \mathbf{x}} =$	$\mathbf{0}^T$ [4]	$\mathbf{0}$ [4]
a is not a function of \mathbf{x} , $u = u(\mathbf{x})$	$\frac{\partial au}{\partial \mathbf{x}} =$		$a \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial(u+v)}{\partial \mathbf{x}} =$		$\frac{\partial u}{\partial \mathbf{x}} + \frac{\partial v}{\partial \mathbf{x}}$
$u = u(\mathbf{x}), v = v(\mathbf{x})$	$\frac{\partial uv}{\partial \mathbf{x}} =$		$u \frac{\partial v}{\partial \mathbf{x}} + v \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x})$	$\frac{\partial g(u)}{\partial \mathbf{x}} =$		$\frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$
$u = u(\mathbf{x})$	$\frac{\partial f(g(u))}{\partial \mathbf{x}} =$		$\frac{\partial f(g)}{\partial g} \frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial \mathbf{x}}$
$\mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial(\mathbf{u} \cdot \mathbf{v})}{\partial \mathbf{x}} = \frac{\partial \mathbf{u}^T \mathbf{v}}{\partial \mathbf{x}} =$	$\mathbf{u}^T \frac{\partial \mathbf{v}}{\partial \mathbf{x}} + \mathbf{v}^T \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$ <ul style="list-style-type: none"> assumes numerator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$ 	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{v} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}} \mathbf{u}$ <ul style="list-style-type: none"> assumes denominator layout of $\frac{\partial \mathbf{u}}{\partial \mathbf{x}}, \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$

Gradient $\nabla E(\mathbf{w})$

$$\nabla E(\mathbf{w}) = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i$$

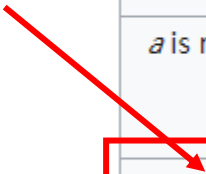
$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}}$$

$$\mathbf{H} = \sum_{i=1}^m \frac{\nabla \{z(\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i\}}{\nabla \mathbf{w}}$$

$$\frac{\nabla z(\mathbf{w}^T \mathbf{x}_i)}{\nabla \mathbf{w}} = z(\mathbf{w}^T \mathbf{x}_i) z(-\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i$$

$$\mathbf{H} = \sum_{i=1}^m \mathbf{x}_i z(\mathbf{w}^T \mathbf{x}_i) z(-\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i^T$$

$a: z(\mathbf{w}^T \mathbf{x}_i)$
 $\mathbf{u}(\mathbf{w}): \mathbf{x}_i$



Identities: vector-by-vector $\frac{\partial \mathbf{y}}{\partial \mathbf{x}}$

Condition	Expression	Numerator layout, i.e. by \mathbf{y} and \mathbf{x}^T	Denominator layout, i.e. by \mathbf{y}^T and \mathbf{x}
\mathbf{a} is not a function of \mathbf{x}	$\frac{\partial \mathbf{a}}{\partial \mathbf{x}} =$	$\mathbf{0}$	
	$\frac{\partial \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{I}	
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{A} \mathbf{x}}{\partial \mathbf{x}} =$	\mathbf{A}	\mathbf{A}^T
\mathbf{A} is not a function of \mathbf{x}	$\frac{\partial \mathbf{x}^T \mathbf{A}}{\partial \mathbf{x}} =$	\mathbf{A}^T	\mathbf{A}
a is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	
$a = a(\mathbf{x}), \mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial a \mathbf{u}}{\partial \mathbf{x}} =$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \mathbf{u} \frac{\partial a}{\partial \mathbf{x}}$	$a \frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial a}{\partial \mathbf{x}} \mathbf{u}^T$
\mathbf{A} is not a function of \mathbf{x} , $\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{A} \mathbf{u}}{\partial \mathbf{x}} =$	$\mathbf{A} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \mathbf{A}^T$
$\mathbf{u} = \mathbf{u}(\mathbf{x}), \mathbf{v} = \mathbf{v}(\mathbf{x})$	$\frac{\partial (\mathbf{u} + \mathbf{v})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} + \frac{\partial \mathbf{v}}{\partial \mathbf{x}}$	
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}}$
$\mathbf{u} = \mathbf{u}(\mathbf{x})$	$\frac{\partial \mathbf{f}(\mathbf{g}(\mathbf{u}))}{\partial \mathbf{x}} =$	$\frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{u}}{\partial \mathbf{x}}$	$\frac{\partial \mathbf{u}}{\partial \mathbf{x}} \frac{\partial \mathbf{g}(\mathbf{u})}{\partial \mathbf{u}} \frac{\partial \mathbf{f}(\mathbf{g})}{\partial \mathbf{g}}$

Gradient $\nabla E(\mathbf{w})$

$$\nabla E(\mathbf{w}) = \sum_{i=1}^m \left[-y_i \mathbf{x}_i + \frac{e^{\mathbf{w}^T \mathbf{x}_i}}{1 + e^{\mathbf{w}^T \mathbf{x}_i}} \mathbf{x}_i \right] = \sum_{i=1}^m [z(\mathbf{w}^T \mathbf{x}_i) - y_i] \mathbf{x}_i = \mathbf{X}^T (\hat{\mathbf{y}} - \mathbf{y})$$

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & x_{12} & \cdots & x_{1d} \\ 1 & x_{21} & x_{22} & \cdots & x_{2d} \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_{m1} & x_{m2} & \cdots & x_{md} \end{pmatrix} = \begin{pmatrix} \mathbf{x}_1^T \\ \mathbf{x}_2^T \\ \vdots \\ \mathbf{x}_m^T \end{pmatrix} \in \mathbb{R}^{m \times (d+1)}, \hat{\mathbf{y}} = \begin{pmatrix} z(\mathbf{w}^T \mathbf{x}_1) \\ z(\mathbf{w}^T \mathbf{x}_2) \\ \vdots \\ z(\mathbf{w}^T \mathbf{x}_m) \end{pmatrix} \in \mathbb{R}^m, \mathbf{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{pmatrix} \in \mathbb{R}^m$$

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}} = \sum_{i=1}^m \mathbf{x}_i z(\mathbf{w}^T \mathbf{x}_i) z(-\mathbf{w}^T \mathbf{x}_i) \mathbf{x}_i^T = \mathbf{X}^T \mathbf{R} \mathbf{X}$$

$\mathbf{R} \in \mathbb{R}^{m \times m}$ is a diagonal matrix with elements $R_{ii} = z(\mathbf{w}^T \mathbf{x}_i) z(-\mathbf{w}^T \mathbf{x}_i)$

- Apply the Newton's method to the logistic regression,

$$\mathbf{w} = \mathbf{w}_t - \mathbf{H}(\mathbf{w}_t)^{-1} \nabla E(\mathbf{w}_t)$$

Compare with Linear Regression

For the linear regression with the sum-of-squares error function, we have,

$$E(\mathbf{w}) = \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 = (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y})$$

$$\nabla E(\mathbf{w}) = \mathbf{X}^T \mathbf{X} \mathbf{w} - \mathbf{X}^T \mathbf{y}$$

$$\mathbf{H} = \nabla^2 E(\mathbf{w}) = \frac{\nabla E(\mathbf{w})}{\nabla \mathbf{w}} = \mathbf{X}^T \mathbf{X}$$

\mathbf{H} is a constant: the error function is quadratic.

Apply the Newton's method to the linear regression,

$$\mathbf{w} = \mathbf{w}_t - \mathbf{H}(\mathbf{w}_t)^{-1} \nabla E(\mathbf{w}_t)$$

$$\mathbf{w} = \mathbf{w}_t - (\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{X} \mathbf{w}_t - \mathbf{X}^T \mathbf{y}) = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \text{Closed-form}$$

The Newton method gives the exact solution in one step.

Summary

Linear Regression

- **Problem**
 - Use hyperplanes to approximate real values
- **Error (Cost) function**
 - Least square
 - $E(\mathbf{w})$: continuous, differentiable, **convex**
- **Algorithm**
 - Analytic solution with pseudo-inverse

Summary

Logistic Regression

➤ Problem

- $P(+1|x)$ as target and $z(\mathbf{w}^T \mathbf{x}_i)$ as hypotheses

➤ Error (Cost) Function

- Negative log-likelihood (cross-entropy)
- $E(\mathbf{w})$: continuous, differentiable, twice-differentiable, **convex**

➤ Optimization

- Iterative methods, e.g., Gradient descent, Newton's method

Exercise

$y \in \{0,1\}$

Target function:

$$f(x) = p(+1|x)$$



$$p(y|x) = \begin{cases} h(x) & \text{for } y = 1 \\ 1 - h(x) & \text{for } y = 0 \end{cases}$$



$y \in \{-1,1\}$

Target function:

$$f(x) = p(+1|x)$$



$$p(y|x) = \begin{cases} h(x) & \text{for } y = 1 \\ 1 - h(x) & \text{for } y = -1 \end{cases}$$

Can you derive the objective function?

Logistic Regression-- $y \in \{-1, 1\}$

Consider $\mathcal{D} = \{(\mathbf{x}_1, +), (\mathbf{x}_2, -), \dots, (\mathbf{x}_m, -)\}$

$$h(\mathbf{x}_i) = P(+1|\mathbf{x}_i) \quad \Leftrightarrow \quad p(y|\mathbf{x}_i) = \begin{cases} h(\mathbf{x}_i) & \text{for } y = +1 \\ 1 - h(\mathbf{x}_i) & \text{for } y = -1 \end{cases}$$

$$\Leftrightarrow p(y|\mathbf{x}_i) = \begin{cases} h(\mathbf{x}_i) & \text{for } y = +1 \\ h(-\mathbf{x}_i) & \text{for } y = -1 \end{cases} \quad \Leftrightarrow \quad p(y|\mathbf{x}_i) = h(y\mathbf{x}_i)$$

$$1 - z(s) = z(-s)$$

$$z(s) = \frac{e^s}{1 + e^s} = \frac{1}{1 + e^{-s}}$$

Logistic Regression-- $y \in \{-1, 1\}$

Consider $\mathcal{D} = \{(\mathbf{x}_1, +), (\mathbf{x}_2, -), \dots, (\mathbf{x}_m, -)\}$

$$\text{likelihood}(h) = \prod_{i=1}^m p(\mathbf{x}_i) p(y_i | \mathbf{x}_i) = p(\mathbf{x}_1) \textcolor{red}{h}(\textcolor{red}{\mathbf{x}}_1) p(\mathbf{x}_2) \textcolor{blue}{h}(\textcolor{blue}{-\mathbf{x}}_2) \dots p(\mathbf{x}_m) \textcolor{blue}{h}(\textcolor{blue}{-\mathbf{x}}_m)$$

$$\max_h \text{likelihood}(h) \propto \prod_{i=1}^m p(y_i | \mathbf{x}_i) = \prod_{i=1}^m h(y_i \mathbf{x}_i) = \prod_{i=1}^m \theta(y_i \mathbf{w}^T \mathbf{x}_i)$$

$$\min_{\textcolor{brown}{w}} - \sum_{i=1}^m \ln \theta(y_i \textcolor{brown}{w}^T \mathbf{x}_i) \quad \Leftrightarrow \quad \min_{\textcolor{brown}{w}} - \sum_{i=1}^m \ln 1 / (1 + e^{-y_i \textcolor{brown}{w}^T \mathbf{x}_i})$$

Cross-entropy loss for
 $y \in \{-1, 1\}$

$$\min_{\textcolor{brown}{w}} - \frac{1 + y_i}{2} \sum_{i=1}^m \ln \frac{1}{1 + e^{-\textcolor{brown}{w}^T \mathbf{x}_i}} - \frac{1 - y_i}{2} \sum_{i=1}^m \ln \frac{1}{1 + e^{\textcolor{brown}{w}^T \mathbf{x}_i}}$$

Logistic Regression-- $y \in \{-1, 1\}$

Cross-entropy

$$H(p, q) = - \sum_x p(x) \log(q(x)) \quad \begin{array}{l} p \in \left\{ \frac{1+y_i}{2}, \frac{1-y_i}{2} \right\} \\ q \in \{h(x), 1-h(x)\} \end{array}$$

$$\min_w - \frac{1+y_i}{2} \sum_{i=1}^m \ln \frac{1}{1+e^{-w^T x_i}} - \frac{1-y_i}{2} \sum_{i=1}^m \ln \frac{1}{1+e^{w^T x_i}}$$

Simplified function

$$\min_w - \sum_{i=1}^m \ln 1/(1 + e^{-y_i w^T x_i})$$