

Machine Learning & Pattern Recognition

Xin Xin(辛鑫)


xinxin@sdu.edu.cn

<https://xinxin-me.github.io/>

Iterative Optimization Technique

Optimization Methods

- Optimization: either minimize or maximize some function $f(x)$ by altering x .
- In most cases, optimization refers to the minimization of $f(x)$.

Maximization $f(x)$  **Minimization** $-f(x)$

- $f(x)$: objective function, cost function, loss function, error function.
- The value that minimize $f(x)$: $x^* = \arg \min f(x)$.

Optimization Methods

- **Deterministic Optimization**
 - The data for the given problem are known accurately.
- **Stochastic Optimization**
 - Refers to a collection of methods for minimizing or maximizing an objective function when randomness is present.

Deterministic Optimization

- First-order methods: methods that use only the gradient.
- Second-order methods: methods that also use the Hessian matrix.

$$\mathbf{H}(f)_{i,j} = \frac{\partial^2}{\partial x_i \partial x_j} f(\mathbf{x})$$

\mathbf{x} : multiple input dimensions.

Newton's Methods

- Motivation: to minimize the local **second-order Taylor** approximation of f .

$$\min_{\mathbf{x}} f(\mathbf{x}) \approx \min_{\mathbf{x}} f(\mathbf{x}_t) + \nabla f(\mathbf{x}_t)^T (\mathbf{x} - \mathbf{x}_t) + \frac{1}{2} (\mathbf{x} - \mathbf{x}_t)^T \nabla^2 f(\mathbf{x}_t) (\mathbf{x} - \mathbf{x}_t)$$

- Take the derivative of \mathbf{x} on both side, we have,

$$\frac{df(\mathbf{x})}{d\mathbf{x}} = \nabla f(\mathbf{x}_t) + \nabla^2 f(\mathbf{x}_t) (\mathbf{x} - \mathbf{x}_t) = \mathbf{0}$$

- Update rule: suppose $\nabla^2 f(\mathbf{x}_t)$ is positive definite,

$$\mathbf{x} = \mathbf{x}_t - [\nabla^2 f(\mathbf{x}_t)]^{-1} \nabla f(\mathbf{x}_t)$$

Newton's Methods

- **Advantage:**

- More **accurate** local approximation of the objective,
- The convergence is much **faster**.

- **Disadvantage:**

- Need to compute the **second derivatives**
- Need to compute the **inverse** of Hessian (time/storage consuming)

Quasi Newton's Methods

- **Main Idea:** To approximate the inverse with a matrix B_t that is iteratively refined by low rank updates to become a better approximation of $[\nabla^2 f(x_t)]^{-1}$.

Quasi Newton's Methods

- BFGS (Broyden–Fletcher–Goldfarb–Shanno):

Broyden, Fletcher, Goldfarb, Shanno



Optimization Methods

- **Deterministic Optimization**
 - The data for the given problem are known accurately.
- **Stochastic Optimization**
 - Refers to a collection of methods for minimizing or maximizing an objective function when randomness is present.

Stochastic Gradient Descent

We now minimize the *empirical risk*,

m is the number of training examples.

$$J(\boldsymbol{\theta}) = \mathbb{E}_{(\mathbf{x}, y) \sim \hat{P}_{data}} L(f(\mathbf{x}, \boldsymbol{\theta}), y) = \frac{1}{m} \sum_{i=1}^m L(f(\mathbf{x}^{(i)}, \boldsymbol{\theta}), y^{(i)})$$

- Optimization algorithms that use the entire training set simultaneously are called *deterministic* or *batch* gradient methods.
- This terminology “*batch*” can be somewhat *confusing*.
 - We use the term “*batch size*” to describe the size of a minibatch in the stochastic gradient descent.
 - We use the term “*batch gradient descent*” to imply the use of the full training set.

Stochastic Gradient Descent

- Optimization algorithms that use only a **single** example at a time are sometimes called **stochastic** or sometimes **online** methods.

- E.g., consider the cost function of linear regression as

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- Now, we ignore the superscript i , then for each x we have ,

$$\begin{aligned} \frac{\partial}{\partial \theta_j} J(\theta) &= \frac{\partial}{\partial \theta_j} \frac{1}{2} (h_{\theta}(x) - y)^2 \\ &= 2 \cdot \frac{1}{2} (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} (h_{\theta}(x) - y) \\ &= (h_{\theta}(x) - y) \cdot \frac{\partial}{\partial \theta_j} \left(\sum_{i=0}^n \theta_i x_i - y \right) \\ &= (h_{\theta}(x) - y) x_j \end{aligned}$$

$$\mathbf{x} \in \mathbb{R}^n$$

Batch vs Stochastic Gradient Descent

- **Batch** gradient descent has to scan through the entire training set before taking a single step—a costly operation if m is large.

Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

λw

Batch GD

- **SGD** can start making progress right away, and continues to make progress with each example it looks at.
- Often, SGD gets θ “close” to the minimum much **faster** than **batch** gradient descent.

Loop {

for i=1 to m, {

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

}

SGD

Batch vs Stochastic Gradient Descent

- SGD may **never “converge”** to the minimum, and the parameters θ will keep oscillating around the minimum of $J(\theta)$;
- **But in practice** most of the values near the minimum will be reasonably **good approximations** to the true minimum.
- Therefore, when the training set is large, SGD is often preferred over batch gradient descent.

Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

Batch GD

Loop {

for i=1 to m, {

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

}

SGD

Stochastic Gradient Descent

- Most algorithms (called *minibatch* or *minibatch stochastic* methods) fall somewhere in between, using more than one but less than all of the training examples.
- **Note:** Confusing again..... It is now common to simply call them *stochastic* methods.

Algorithm 8.1 Stochastic gradient descent (SGD) update at training iteration k

Require: Learning rate ϵ_k .

Require: Initial parameter θ

while stopping criterion not met **do**

 Sample a minibatch of m examples from the training set $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ with corresponding targets $\mathbf{y}^{(i)}$.

 Compute gradient estimate: $\hat{\mathbf{g}} \leftarrow +\frac{1}{m} \nabla_{\theta} \sum_i L(f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)})$

 Apply update: $\theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$


end while

Accelerated SGD for Deep Learning


- **Polyak's Classical Momentum** [Polyak 1964]
- **Nesterov's Momentum** [Nesterov 1983]

Polyak's Classical Momentum [Polyak 1964]

- **Motivation:** Key problem of Gradient Descent is “zig-zagging”.



(a) Zig-zagging problem of GD.



(a) Trajectory of CM on same problem.

Polyak's Classical Momentum [Polyak 1964]

- The classical momentum (CM) accumulates an **exponentially decaying moving average of past gradients** and continues to move in their direction.
- Letting η be the learning rate.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$

$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t} f$$

- Velocity vector \mathbf{v}_t : a **memory** that accumulates the directions of reduction that were chosen in the previous t steps.
- The influence of \mathbf{v} is controlled by the **momentum coefficient** $\mu \in [0,1]$. μ is usually slightly less than 1. When $\mu = 0$:

Polyak's Classical Momentum [Polyak 1964]

- The classical momentum (CM) accumulates an **exponentially decaying moving average of past gradients** and continues to move in their direction.
- Letting η be the learning rate.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$

$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t} f$$

- Velocity vector \mathbf{v}_t : a **memory** that accumulates the directions of reduction that were chosen in the previous t steps.
- The influence of \mathbf{v} is controlled by the **momentum coefficient** $\mu \in [0,1]$. μ is usually slightly less than 1. When $\mu = 0$: it is just the **Gradient Descent**.

Polyak's Classical Momentum [Polyak 1964]

- Exponential decay

$$\begin{aligned}\mathbf{v}_{t+1} &= \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t} f = \mu(\mu \mathbf{v}_{t-1} - \eta \nabla_{\mathbf{w}_{t-1}} f) - \eta \nabla_{\mathbf{w}_t} f \\ &= \mu^{t+1} \mathbf{v}_0 - (\mu^t \eta \nabla_{\mathbf{w}_0} f + \cdots + \mu^1 \eta \nabla_{\mathbf{w}_{t-1}} f + \eta \nabla_{\mathbf{w}_t} f)\end{aligned}$$

We know that $\lim_{n \rightarrow \infty} (1 + \frac{1}{n})^n = e$ $\lim_{n \rightarrow \infty} (1 - \frac{1}{n})^n = \frac{1}{e}$

Let $x = \frac{1}{n} \rightarrow 0$, we have $(1 - x)^{\frac{1}{x}} = \frac{1}{e}$

We 'ignore' the terms whose weights decay to less than $\frac{1}{e}$.

For example, if $\mu = 0.9$ ($x = 0.1$), to calculate \mathbf{v}_{100} , we only consider the last 10 steps (i.e., $\mathbf{v}_{99} \dots \mathbf{v}_{90}$) as the valid memory.

Polyak's Classical Momentum [Polyak 1964]

- The SGD algorithm with **momentum** is given as follows.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1} \quad \mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t} f$$

Algorithm 8.2 Stochastic gradient descent (SGD) with momentum

Require: Learning rate η , momentum parameter μ .

Require: Initial parameter $\boldsymbol{\theta}$, initial velocity \mathbf{v} .

while stopping criterion not met **do**

 Sample a minibatch of m examples from the training set $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ with corresponding targets $\mathbf{y}^{(i)}$.


 Compute gradient estimate: $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\boldsymbol{\theta}} \sum_i L(f(\mathbf{x}^{(i)}; \boldsymbol{\theta}), \mathbf{y}^{(i)})$

 Compute velocity update: $\mathbf{v} \leftarrow \mu \mathbf{v} - \eta \mathbf{g}$

 Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{v}$

end while

Polyak's Classical Momentum [Polyak 1964]




Nesterov's Accelerated Gradient [Nesterov 1983]

- The update equations of NAG are:

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$

$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t + \mu \mathbf{v}_t} f$$




NAG

CM

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$

$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t} f$$



CM

Illustration of the comparison between CM and NAG.

Nesterov's Accelerated Gradient [Nesterov 1983]

- **First** make a big jump in the direction of the previous accumulated gradient.
- **Then** measure the gradient where you end up and make a **correction**.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$


$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\boxed{\mathbf{w}_t + \mu \mathbf{v}_t}} f$$

NAG

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1}$$

$$\mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\boxed{\mathbf{w}_t}} f$$

CM



brown vector = jump, red vector = correction, green vector = accumulated gradient
blue vectors = standard momentum

Nesterov's Accelerated Gradient [Nesterov 1983]

- The update equations of NAG are:

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t + \mathbf{v}_{t+1} \\ \mathbf{v}_{t+1} &= \mu \mathbf{v}_t - \eta \nabla_{\boxed{\mathbf{w}_t + \mu \mathbf{v}_t}} f \end{aligned}$$

CM

$$\begin{aligned} \mathbf{w}_{t+1} &= \mathbf{w}_t + \mathbf{v}_{t+1} \\ \mathbf{v}_{t+1} &= \mu \mathbf{v}_t - \eta \nabla_{\boxed{\mathbf{w}_t}} f \end{aligned}$$

- CM → inspecting the gradient at the **current iterate** of \mathbf{w}_t ;
 - Faithfully trusts the current iterate;
- NAG → inspecting **the gradient at $\mathbf{w}_t + \mu \mathbf{v}_t$** .
 - Puts less faith into the current iterate and **looks ahead** in the direction suggested by the velocity vector.
- The small difference allows NAG to adapt faster and in a more stable way.

Nesterov's Accelerated Gradient [Nesterov 1983]

- The complete Nesterov momentum algorithm is presented as follows.

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \mathbf{v}_{t+1} \quad \mathbf{v}_{t+1} = \mu \mathbf{v}_t - \eta \nabla_{\mathbf{w}_t + \mu \mathbf{v}_t} f$$

Algorithm 8.3 Stochastic gradient descent (SGD) with Nesterov momentum

Require: Learning rate η , momentum parameter μ .

Require: Initial parameter $\boldsymbol{\theta}$, initial velocity \mathbf{v} .

while stopping criterion not met **do**

 Sample a minibatch of m examples from the training set $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ with corresponding labels $\mathbf{y}^{(i)}$.

 Apply interim update: $\tilde{\boldsymbol{\theta}} \leftarrow \boldsymbol{\theta} + \mu \mathbf{v}$

 Compute gradient (at interim point): $\mathbf{g} \leftarrow \frac{1}{m} \nabla_{\tilde{\boldsymbol{\theta}}} \sum_i L(f(\mathbf{x}^{(i)}; \tilde{\boldsymbol{\theta}}), \mathbf{y}^{(i)})$

 Compute velocity update: $\mathbf{v} \leftarrow \mu \mathbf{v} - \eta \mathbf{g}$

 Apply update: $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{v}$

end while

Algorithms with Adaptive Learning Rates

- Researchers have long realized that the learning rate is the most difficult hyperparameter to set because it has a significant impact on model performance.
- Recently, a number of methods have been introduced that adapt the learning rates of model parameters.
 - Adagrad
 - RMSProp
 - Adam

Adagrad

Adagrad (Duchi et al, COLT 2010) uses a different learning rate for **every parameter** θ_i at **every time step** t . It individually adapts the learning rates of all model parameters by scaling them inversely proportional to the square root of the sum of all of their historical squared values

$$g_{t,i} = \nabla_{\theta_t} J(\theta_{t,i})$$

The SGD update

$$\theta_{t+1,i} = \theta_{t,i} - \eta \cdot g_{t,i}$$

$$\theta_{t+1,i} = \theta_{t,i} - \frac{\eta}{\sqrt{G_{t,ii} + \epsilon}} \cdot g_{t,i}$$

$G_t \in \mathbb{R}^{d \times d}$ is a diagonal matrix where each diagonal element i, i is **the sum** of the **squares** of the **gradients** w.r.t. θ_i **up to time step t** .

Weakness: Since every added term is positive, the accumulated sum **keeps growing** which in turn causes the learning rate to **shrink** and eventually become infinitely small.

Adadelta & RMSprop

Adadelta is an extension of Adagrad that seeks to reduce its aggressive, monotonically decreasing learning rate. Instead of accumulating all past squared gradients, Adadelta **restricts the window of** accumulated past gradients to some **fixed size w** .

Via a **decaying** mechanism, $E[g^2]_t = \gamma E[g^2]_{t-1} + (1 - \gamma)g_t^2$

RMSprop is an unpublished, adaptive learning rate method proposed by Geoff Hinton in his Coursera Class. RMSprop and Adadelta have both been developed **independently** around the same time to resolve Adagrad's radically diminishing learning rates.

衰减速率：可调超参

$$E[g^2]_t = 0.9E[g^2]_{t-1} + 0.1g_t^2$$
$$\theta_{t+1} = \theta_t - \frac{\eta}{\sqrt{E[g^2]_t + \epsilon}}g_t$$

采用指数衰减平均，以丢弃遥远过去的历史。

Adam

- **Adaptive Moment Estimation (Adam)** is another method that computes adaptive learning rates for each parameter.
- Adam was presented by Diederik Kingma from OpenAI and Jimmy Ba from the University of Toronto in their **2015** ICLR paper (poster) titled “Adam: A Method for Stochastic Optimization”.
 1. Adam stores an **exponentially decaying** average of past **squared gradients** \mathbf{v}_t (variance) like **Adadelta** and **RMSprop**.
 2. Adam also keeps an **exponentially decaying** average of **past gradients** \mathbf{m}_t (mean), similar to **momentum**.

Adam

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

✓ $v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

As m_t, v_t are initialized as 0, they are biased towards 0, especially when $\beta_1 \rightarrow 1, \beta_2 \rightarrow 1$. →

8.5.4 Choosing the Right Optimization Algorithm

In this section, we discussed a series of related algorithms that each seek to address the challenge of optimizing deep models by adapting the learning rate for each model parameter. At this point, a natural question is: which algorithm should one choose?

Unfortunately, there is currently no consensus on this point. Schaul *et al.* (2014) presented a valuable comparison of a large number of optimization algorithms across a wide range of learning tasks. While the results suggest that the family of algorithms with adaptive learning rates (represented by RMSProp and AdaDelta) performed fairly robustly, no single best algorithm has emerged.

Currently, the most popular optimization algorithms actively in use include SGD, SGD with momentum, RMSProp, RMSProp with momentum, AdaDelta and Adam. The choice of which algorithm to use, at this point, seems to depend largely on the user's familiarity with the algorithm (for ease of hyperparameter tuning).

Summary

- **Deterministic Optimization**
 - The data for the given problem are known accurately.
 - **First** order method: e.g., Gradient Descent.
 - **Second** order method: e.g., **Newton's** method, Quasi Newton's Methods (BFGS).
- **Stochastic Optimization**
 - Using several samples of the training examples (minibatch).
 - SGD
 - SGD+Acceleration
 - Polyak's Classical Momentum
 - **Nesterov's Momentum**
 - Agrad/Adadelata/RMSprop/**Adam**